Ishan Rajvi

AI/ML ENGINEER

Location: Fort Worth, TX 76114 | Email: ishandrajvi1904@gmail.com | Phone: (817) 947-4436

https://www.linkedin.com/in/ishan-rajvi/

SUMMARY

AI/ML Engineer with 4+ years of experience building and deploying production-grade machine learning systems across enterprise environments. Proven track record of designing end-to-end ML pipelines from data ingestion and feature engineering to model deployment, monitoring, and optimization. Specialized in NLP, predictive modeling, and Generative AI (LLMs, RAG), with hands-on experience delivering cloud-native solutions on AWS and Azure that drive automation, reliability, and measurable business impact.

SKILLS

Languages:	Python, R, SQL
Libraries & Frameworks:	TensorFlow, PyTorch, Keras, Scikit-learn, Pandas, NumPy, Matplotlib, Seaborn, BeautifulSoup, Hugging Face, LangChain, LangGraph, NLTK, spaCy, SciSpaCy
Generative AI & LLMs:	GPT-4, Claude, BERT, RoBERTa, Transformers, RAG, LoRA, QLoRA, PEFT, CrewAI
Vector Databases:	FAISS, Pinecone, Weaviate
Machine Learning:	Classification, Regression, Clustering, Time-Series Analysis, Reinforcement Learning, Transfer Learning, GANs
Data Engineering:	PySpark, Airflow, Databricks, ETL/ELT Pipelines, Data Modeling, Hadoop, Teradata, BigQuery, Snowflake
Cloud Platforms:	AWS, Azure, IBM Cloud (AutoAI, Watson Assistant)
MLOps & Deployment:	MLflow, Docker, Kubernetes, Kubeflow, CI/CD Pipelines, REST APIs (Flask, FastAPI)
Visualization & BI:	Power BI, Tableau, Looker
Security & Governance:	Responsible AI, Model Explainability (SHAP, LIME), Bias Mitigation, Privacy-aware ML
Other:	Data Structures & Algorithms, Advanced Excel, Agile/Scrum, Jira, SDLC

EXPERIENCE

Pfizer, USA **| Aug 2025 – Present**

AI/ML Engineer

- Led the design and production deployment of predictive ML systems on Azure Data Lake telemetry and machine logs, enabling proactive failure detection and reducing incident recurrence by 18% across critical network infrastructure.
- Architected an NLP-driven ticket automation platform that classified and routed IT incidents at scale, cutting manual triage effort by 40% and accelerating mean-time-to-resolution (MTTR).
- Designed custom text preprocessing and representation pipelines (tokenization, lemmatization, domain-specific vocabularies) to improve precision/recall of production NLP models.
- Partnered with operations and leadership teams to deliver real-time Power BI dashboards integrating predictive risk signals with operational KPIs, directly supporting executive decision-making.

Hexaware Technologies, India **| Jun 2022 – Jul 2023**

AI/ML Engineer

- Designed and deployed production-grade ML systems across NLP, computer vision, and recommendation use cases, supporting large-scale enterprise applications.
- Built automated data pipelines using Python, SQL, and Spark, ensuring high-quality, reproducible datasets for model training and inference.
- Implemented full MLOps infrastructure using MLflow, Airflow, Docker, and Kubernetes, enabling automated training, versioning, monitoring, and rollback of models in production.
- Built end-to-end feature engineering, model training, validation, and monitoring pipelines, owning the full ML lifecycle from raw data to business-facing insights.
- Optimized models through hyperparameter tuning and cross-validation, improving prediction accuracy while reducing inference latency in real-time systems.
- Developed RESTful microservices with Flask and FastAPI for scalable model serving.
- Worked closely with product, engineering, and business teams to translate requirements into deployed AI solutions, reducing production model drift by 20%.

Karbh IT Solutions, India **| Mar 2021 – May 2022**

Data Scientist

- Developed and evaluated multiple machine learning models (XGBoost, Random Forest, SVM, Logistic Regression) for enterprise classification and forecasting problems, achieving up to 15% performance gains over baseline systems.
- Applied Bayesian Hidden Markov Models for data analysis, improving forecast stability and model interpretability in time-series use cases.
- Conducted large-scale EDA and feature engineering using Pandas, NumPy, and SciPy to uncover business-critical patterns and data quality issues.
- Built NLP pipelines for text cleaning and sentiment analysis, improving model performance on unstructured customer data.
- Designed Tableau dashboards that reduced reporting turnaround time and improved stakeholder visibility into key business metrics.
- Integrated structured and unstructured data from SQL, MongoDB, Hadoop, and web sources to support end-to-end ML workflows.

EDUCATION

M.S. in Computer Science | *The University of Texas, Arlington, TX, USA* **| May 2025**

B.E. in Computer Engineering | *L.D. College of Engineering, India* **| Jun 2022**